Benchmarking Counterfactual Explanations of Model Change with Exact Ground Truth

Bao B. Dinh, Minh C. Tran and Bac Le

Faculty of Information Technology, University of Science
Vietnam National University, Ho Chi Minh City - Vietnam

Abstract. Methods that explain a model update are rarely checked against ground truth. We build a benchmark with exactly known changes, null updates included, and give a linear analysis of what five families of explanations detect, among them differential counterfactuals (Δ-CE): the nearest inputs on which only one model changes its decision, both decisively. As predicted, gradient differences miss threshold shifts and detect rotations on small CNNs, though not on a ResNet-8. Δ-CE stays silent on gently fine-tuned null updates, where counterfactual and attribution differences report change, and localises two independent changes, but a pre-specified hypothesis that it names the changed factor more often than gradient differences was rejected. Its silence comes from the margin check; a post hoc analysis shows that the search adds witnesses that a filtered ordinary counterfactual misses. Our results suggest that explanations of change need null updates as a reference. Code and supplement: *github.com/Johnlrr/explaining-model-change*.

1 Introduction

An update that improves accuracy can still change individual decisions in unwanted ways [1]. Explaining *how* the models differ is a goal of explainable AI under distributional change [2, 3], distinct from explaining a shift of the data [4]. Artelt et al. [5, 6] compare the counterfactuals [7] of one input under each model, or take a difference of gradients; Muschalik et al. [8] compare feature importances; disagreement search [9, 10, 11] looks for inputs on which the models disagree. Ground-truth benchmarks exist for explaining one model [12, 13, 14], not a model update, so what each family can see goes unchecked.

We say that a method *sees* a change if it stays silent when nothing changed (abstention), reports a change when there is one (*whether*), and names the changed factor and region (*what*, *where*). As a fifth family we study *differential counterfactuals* (Δ-CE): the closest on-manifold input where only f_{new} changes its decision, with both models deciding by a margin of at least κ. Unlike robust counterfactuals [15, 16], valid for both models, it is valid for exactly one.

Contributions. (i) A benchmark with exact ground truth: five rendered factors, four shift types, null updates and a 4×4 grid over rotation and threshold. (ii) A linear analysis separating rotation from offset, tested on CNNs and a ResNet-8 under a protocol fixed in advance; all failures are reported. (iii) Δ-CE stays silent on gently fine-tuned null updates, also unscreened and through a learned decoder, where counterfactual and attribution differences report change; it localises two independent changes and finds witnesses that filtering misses.

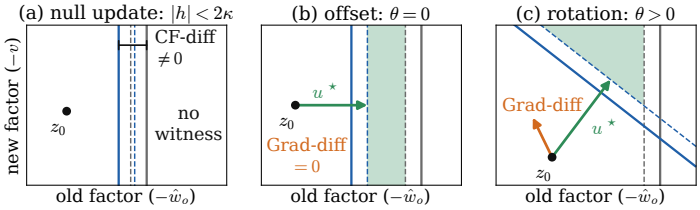


Fig. 1: Linear case. Solid: boundaries of f_{old} (grey), f_{new} (blue); dashed: $m_o = \kappa$, $m_n = -\kappa$; shaded: witnesses. (a) None exists (Lemma 1). (b), (c) Gradient difference and witness $u^\star$ (Proposition 2).

2 Differential counterfactuals and what each family sees

Let $D : \mathcal{Z} \to \mathcal{X}$ map a search space onto the data manifold (here an exact renderer from factor space $\mathcal{Z} = [0,1]^5$). For a query z_0 that both models assign to class y, the margin is the logit difference $m(z) = f_y(D(z)) - f_{1-y}(D(z))$; m_o, m_n belong to f_{old}, f_{new}, and $h = m_n - m_o$; z_0 is valid if $m_o(z_0), m_n(z_0) \geq \kappa$.

Definition 1. With $\mathcal{F}_\kappa^N = \{z : m_n(z) \leq -\kappa,\ m_o(z) \geq \kappa\}$, a Δ-CE of z_0 is $z^\star \in \arg\min_{z \in \mathcal{F}_\kappa^N} \|z - z_0\|_2$ if $\|z^\star - z_0\|_2 \leq r$, and NOT_FOUND otherwise; swapping the models gives type O. Each $z \in \mathcal{F}_\kappa^N$ is a *witness*: two forward passes verify that the models disagree, both decisively.

Lemma 1. Every witness has $h(z) \leq -\kappa - m_o(z) \leq -2\kappa$, so models whose margins differ by less than 2κ everywhere have none.

Corollary 1 (constant logit offset). If f_{new} adds b to the logit of class 1, e.g. the prior correction [17], then $h \equiv \pm b$ and $\nabla h \equiv 0$. Type-N witnesses exist only for the class whose margin b lowers and only if $|b| \geq 2\kappa$, at the points with $\kappa \leq m_o \leq |b| - \kappa$. *Proof:* insert $h \equiv -|b|$ in Lemma 1; no linearity is needed.

Linear case (Fig. 1)**.** For linear margins, let the unit normals be $\hat{w}_n = \cos\theta\, \hat{w}_o + \sin\theta\, v$, $v \perp \hat{w}_o$ ($\theta = 0$: offset; $\theta > 0$: rotation towards v). Let d_n, d_o be the distances from z_0 to $\{m_n = -\kappa\}$ and $\{m_o = \kappa\}$, and u_n, u_o the steps to the nearest counterfactuals of f_{new}, f_{old}. We compare the gradient difference [5] $-(\hat{w}_n - \hat{w}_o)$, counterfactual differencing $u_n - u_o$ and the witness $u^\star = z^\star - z_0$.

Proposition 2. (a) *Offset* ($\theta = 0$): the gradient difference is 0; counterfactual differencing is a multiple of $\hat{w}_o$; a witness exists iff $d_n \leq d_o$ and equals $-d_n \hat{w}_o$. (b) *Rotation* ($0 < \theta < \pi/2$): the gradient difference has a larger component along v than along $\hat{w}_o$ for every θ, and so does the witness iff $\theta > \pi/4$ or $d_o(\sin\theta + \cos\theta) < d_n$. After projecting out $\hat{w}_o$, all three coincide up to positive scale, as in Theorem 1 of [5] (new here: the offset case and the witness). *Proof* of (a): both normals equal $\hat{w}_o$; of (b): online supplement.

3 Benchmark and protocol

Data and models. Images are 32×32 ellipses rendered from five independent factors in $[0,1]$: aspect (task factor), size, colour, background and horizontal position; labels are $y = \mathbb{1}[\text{aspect} > 0.5]$ with 5% noise. The exact renderer gives

Table 1: Adaptation conditions; $k^\star$: ground-truth factor.

Change	Shift	Adaptation data	$k^\star$	Pairs
Shortcut	spurious	colour correlated with label (0.99)	colour	3
Subgroup	local concept	labels flipped for size > 0.75	size	3
Threshold	concept	$y = \mathbb{1}[\text{aspect} > 0.6]$	aspect	3
Null	none / covariate	resample / size $\sim U[0.5, 1]$, same rule	—	5 / 3
Grid[a]	concept	$y = \mathbb{1}[\text{aspect} + \alpha(\text{colour} - 0.5) > t]$	aspect / colour	8 / 24
Offsets[b]	concept	grid, $\alpha = 0$, $t \in \{.53, .56, .6\}$, seeds 2–4	aspect	9
Both[c]	spurious + local	Shortcut and Subgroup at once	colour, size	8

[a] $\alpha \in \{0, .1, .25, .5\} \times t \in \{.5, .53, .56, .6\}$, 2 seeds; 4 offsets tested. [b] 8 qualified. [b,c] Registered.

the factor change of any explanation without a learned oracle, and all methods search the same factor space through it. f_{old} is a small CNN (five seeds); f_{new} is f_{old} fine-tuned on 10 000 samples from one condition (Table 1). In the sense of Section 2, a higher label threshold is an offset and a rule that adds colour is a rotation. A pair enters only if its change alters validation decisions. The ground truth is the injected change; what f_{new} learns is read from the realised rotation $\hat{\alpha}$ and threshold $\hat{t}$ of a logistic fit to its decisions, at which the grid is analysed.

Metrics and methods. On $N = 300$ valid test points per pair, Recovery@r is the share with a valid explanation within the budget r whose largest factor change is $k^\star$ (NOT_FOUND and ties count 0); the found-rate is the share with any valid explanation within r. For Δ-CE we count type-N witnesses only (pre-declared), so a query whose class the update raises counts 0. All counterfactual methods share one solver, the test points and the budget $r = 0.54$, set once on validation. A margin score, $|h(z_0)|$, serves as the control for *whether*. We compare Δ-CE ($\kappa = 0.5$) with **CF-diff**, counterfactual differencing [5]; **Grad-diff**, the gradient difference [5]; **Imp-diff**, a difference of permutation importances [8], one vector per pair; **Seg-search**, a line search towards training points, after [9]; **No-prox**, Δ-CE without proximity, the DeepXplore objective [10], to which DiG-IN [11] corresponds on an exact manifold; **CF-new**, the counterfactual of f_{new}; and, in Section 4.2, IG- [18] and SHAP-differences [19].

Protocol. Each test set served at most one confirmatory test, after all choices were fixed. The main test set carries Table 2; hypotheses formed after it were tested on new data, and three follow-ups were registered before their data existed [20]: unscreened nulls, offsets and a ResNet-8; the learned decoder; and two changes at once. Bootstrap intervals resample points within pairs, so Table 3 also counts pairs; the supplement repeats each test with the pair as unit and dates all test sets. A digit-image task and a label-shift update failed qualification.

4 Results

4.1 Abstention, detection and localisation

Abstention. No null query has a witness: on the 8 null pairs of Table 2, five null updates of an MLP on Adult, and 8 registered unscreened null pairs, admitted without the qualification check. On the same nulls (covariate-shifted included) CF-diff reports a change for 97.7–100% of points. Silence does not need the exact

Table 2: Recovery on the main test set (↑) and false reports on null updates (↓).

Condition	Δ-CE	CF-diff	Grad-diff	Imp-diff	Seg-search	No-prox	CF-new
Shortcut ↑	0.38 (0.84)	0.54 (0.93)	**0.91** (1.0)	0.00	0.02 (0.08)	0.35 (0.63)	0.07 (0.93)
Subgroup ↑	0.53 (0.81)	0.32 (0.99)	**0.62** (1.0)	0.00	0.06 (0.10)	0.29 (0.58)	0.30 (1.0)
Threshold ↑	0.47 (0.47)	0.76 (0.90)	0.07 (1.0)	**1.00**[a]	0.07 (0.08)	0.31 (0.40)	0.90 (0.90)
Null ↓	**0.000**	0.994	1.0	1.0	**0.000**	**0.000**	0.996

Change rows: Recovery@r (found-rate), mean of 3 pairs, CIs ±0.035. Null: found-rate on 8 null pairs; any output is a false report. Bold: best in row. [a]Ranks aspect first in every pair.

Table 3: Tests and their outcomes (✓ confirmed, × not confirmed).

Claim	Result	Data	Estimate	p	Pairs
Consistency checks (expected from Section 2)					
No witness on nulls	✓ post	Nulls; Adult MLP	0/2 400; 0/1 500	–	8/8; 5/5
CF-diff flags $\|b\| < 2\kappa$	✓ post	Logit offsets, uncut	0.961 [0.955, 0.967]	$2{\cdot}10^{-4}$	–
Δ-CE > Grad-diff	(✓) post	Threshold	+0.31 [0.27, 0.35]	$2{\cdot}10^{-4}$	3/3
Abstention, detection and localisation					
No witness on nulls	✓ OSF	Unscreened; decoder	0/2 400; 0/800	–	8/8; 8/8
No witness on nulls	× OSF; expl.	ResNet-8, $\kappa = 0.5$; 2.5	180/200; 0/200	–	0/2; 2/2
Witness ≥ margin score +0.2	× OSF	Registered offsets	+0.17 to +0.97	–	7/8
Witness localises ($L > 0$)	✓ OSF	Two changes	$L = 0.50$	0.004	8/8
Offsets					
Grad-diff Recovery < 0.2	(✓) pre	Grid, $\alpha = 0$	0.081 [0.066, 0.096]	$4{\cdot}10^{-4}$	4/4
Grad-diff Recovery < 0.2	✓ OSF	Registered offsets	0.033–0.093	0.004	8/8
Grad-diff offset-invariant	× pre	Grid	−0.098 [−0.143, −0.053]	0.94	–
Rotations					
Grad-diff rises with α	✓ pre	Grid	+0.57 [0.53, 0.61]	$4{\cdot}10^{-4}$	$\rho = 0.89$
Grad-diff < 0.2; rot. > offset	× OSF	ResNet-8	0.32–0.36; 0.06–0.10	–	0/3; 0/3
Projected Δ-CE ≈ Grad-diff	× post; pre	Shortcut; grid	−0.30; −0.35 (±0.03)	1; 1	–
Normal at $z^\star$ halves the gap	(✓) post	Grid	−0.19 → +0.03 [0.01, 0.05]	$6{\cdot}10^{-4}$	4/4
Δ-CE versus gradient differences					
Δ-CE > Grad-diff (rejected)	× pre	Subgroup	−0.09 [−0.13, −0.05]	–	0/3
Δ-CE ≈ Grad-diff, CF-diff	× pre	Shortc., Thresh.	$\|\Delta\| = 0.17$–0.53	–	–

(✓): confirmed over points only (≤ 4 pairs). When the test was fixed: OSF, registered [20] before its data existed; pre, before any result or, for the grid, before its run; post, after earlier results, then tested on new data; expl., exploratory. Estimate: witnesses on nulls, Recovery, found-rate or a Recovery difference Δ (≈, invariant: $|\Delta| \le 0.1$); 95% CIs (90% for band and halving tests). p: sign test over pairs (OSF rows) or Holm-adjusted bootstrap (families in the supplement). Pairs: on the claimed side / all; ρ: rank correlation over 24 pairs.

renderer: searching through the decoder of an autoencoder trained on the images (registered), Δ-CE finds no witness for a null query (0/800; CF-diff 85–94%), and on 6 of 8 offsets finds one for at least half of the lowered-class queries.

With $\kappa = 0$, witnesses appear on 83% of null queries. Silence also depends on the recipe: the nulls are fine-tuned more gently than the offsets (fewer epochs, lower learning rate; supplement), though their threshold moves as far ($|\Delta\hat{t}| \le 0.035$); the grid's null cell, fine-tuned like the offsets, has witnesses on 22% of queries at $\kappa = 0.5$ and none at $\kappa = 1$, where small offsets keep few witnesses (0.04). Silence is thus a property of κ relative to the size of the update, as Lemma 1 states, and not of the search.

Detection. With each score cut on the nulls to flag none, witnesses exist for 0.17–1.00 of lowered-class queries on the 8 registered offsets, the margin score flags 0.00–0.54 and CF-diff none. The witness beat the margin score by 0.2 in 7 of 8 pairs, short of the registered 8/8, so we claim no advantage for *whether*.

Localisation. With a colour shortcut everywhere and a label flip for large objects, witnesses change size most for 46–75% of queries near the large-object

region and 13–23% far from it; this difference L is positive in all 8 pairs (pooled 0.50; Table 3). Grad-diff and CF-diff barely localise ($L = 0.09$ and 0.13), and Imp-diff, one vector per pair, cannot; CF-new localises equally well (0.51).

Search versus filtering (post hoc, saved outputs). The counterfactual of f_{new} kept only where $m_o \geq \kappa$ (filtered CF-new) is a witness; it is silent on the nulls and localises ($L = 0.48$), but finds witnesses for only 0.42–0.46 of the main-test queries, against 0.47–0.84 for Δ-CE. The extra witnesses carry Δ-CE's Recovery on Shortcut (0.38 vs. 0.03) and part on Subgroup (0.53 vs. 0.30), higher in 3 of 3 pairs each, but not on Threshold (0.47 vs. 0.45). Null-adjusted (minus the rate of naming $k^\star$ on nulls), Δ-CE leads on Subgroup and Threshold, Grad-diff on Shortcut (0.79); unfiltered CF-new falls to −0.09 on Threshold. The search costs 2.4 s per query (CF-new: 0.6 s; Grad-diff: under 1 ms).

4.2 Offsets, rotations and further baselines

Gradient differences are blind to an offset, as Proposition 2(a) predicts: their Recovery is 0.081 on the grid's pure offsets and 0.03–0.09 on the 8 registered offsets (Table 3). On Threshold, Δ-CE beats them on new data (0.41 vs. 0.10; main test: 0.47 vs. 0.07) in 3 of 3 pairs, too few for a pair-level test. Type-N witnesses exist only for the lowered class (class 1), hence Δ-CE's found-rate of 0.47; type-O witnesses for the other class raise Recovery to 1.00 on new data.

Gradient differences see a rotation, increasingly with its angle: Recovery rises from 0.36 at $\alpha = 0.1$ to 0.93 at $\alpha = 0.5$ ($t = 0.5$; Spearman $\rho = 0.89$ over the 24 rotation pairs; map of all pairs in the supplement). The projection equivalence of Proposition 2(b) holds only with $\hat{w}_o$ taken at the witness (Table 3).

Attribution differences and other baselines. With predictions fixed beforehand, IG- and SHAP-diff recover colour on Shortcut (0.93) and, against prediction, name aspect on Threshold (0.73–0.75; Grad-diff 0.07), but also on 57% and 71% of null queries (IG-diff on 42% of the registered nulls, as predicted), so this carries no signal. Imp-diff ranks aspect first in all 9 change pairs on test data, but ranks $k^\star$ first on adaptation data. Seg-search with 10× the segments finds witnesses for 0.26–0.47 of queries (predicted ≤ 0.2), none missed by Δ-CE.

Second architecture. On a registered ResNet-8 (run locally after its cost rule had excluded it), null fine-tuning moves margins four times more than on the CNNs, and witnesses appear on 0.90 of null queries: κ is model-specific. With κ calibrated on validation nulls (rule fixed in advance; grid extended after $\kappa \leq 2$ failed), giving $\kappa^\star = 2.5$, no new null query has a witness (0/200) and witnesses exist for 0.78–1.00 of lowered-class offset queries (Δ-CE Recovery 0.36–0.40). Grad-diff sees neither change: it names aspect on 0.32–0.36 of offset queries, near its 0.20–0.30 on nulls, and its rotation Recovery, 0.06–0.10, is below that of the offset with equal seed and threshold in all 3 pairs.

5 Discussion and limitations

Our results suggest that an explanation of change needs null updates as a reference, since counterfactual and attribution differences report change on them;

on small CNNs, gradient differences also miss offsets. Abstention comes from the check $m_o \geq \kappa$, which can also filter the counterfactual of f_{new}; what is specific to Δ-CE is the search for witnesses at queries where that counterfactual fails the check. **Limitations.** (1) The data are synthetic and binary; Recovery through the learned decoder is approximate and lower for Δ-CE than for CF-diff; multi-class tasks, other domains, retraining from scratch and architecture changes are untested, and we did not test why the linear prediction fails on the ResNet-8; the tabular test covers abstention only. (2) κ is in logit units, exact only for a constant offset, and must be calibrated on null updates of each model family and fine-tuning recipe, obtained by fine-tuning f_{old} on resampled data. (3) Witnesses see only decisive changes on the manifold within the budget, without a certificate. (4) Several tests were formed after earlier results; changes were logged before their data existed, with three disclosed exceptions.

References

[1] S. Yan et al. Positive-congruent training: towards regression-free model updates. *CVPR*, 2021.
[2] B. Hammer, E. Hüllermeier. Interpretable machine learning: on the problem of explaining model change. *Workshop Comput. Intell.*, 2021.
[3] M. Kaden et al. Reliable counterfactuals for machine learning models – current aspects and perspectives. *ESANN*, 2026.
[4] S. Kulinski, D. I. Inouye. Towards explaining distribution shifts. *ICML*, 2023.
[5] A. Artelt et al. Contrastive explanations for explaining model adaptations. arXiv, 2021.
[6] A. Artelt et al. Contrasting explanations for understanding and regularizing model adaptations. *Neural Process. Lett.*, 2022.
[7] S. Wachter et al. Counterfactual explanations without opening the black box. *Harv. JOLT*, 2017.
[8] M. Muschalik et al. Agnostic explanation of model change based on feature importance. *KI*, 2022.
[9] X. Renard et al. Understanding prediction discrepancies in machine learning classifiers. arXiv, 2021.
[10] K. Pei et al. DeepXplore: automated whitebox testing of deep learning systems. *SOSP*, 2017.
[11] M. Augustin et al. DiG-IN: diffusion guidance for investigating networks. *CVPR*, 2024.
[12] C. Agarwal et al. OpenXAI: towards a transparent evaluation of model explanations. *NeurIPS*, 2022.
[13] R. Hesse et al. FunnyBirds: a synthetic vision dataset for a part-based analysis of explainable AI methods. *ICCV*, 2023.
[14] L. Arras et al. CLEVR-XAI: a benchmark dataset for the ground truth evaluation of neural network explanations. *Inf. Fusion*, 2022.
[15] F. Hamman et al. Robust counterfactual explanations for neural networks with probabilistic guarantees. *ICML*, 2023.
[16] M. Pawelczyk et al. On counterfactual explanations under predictive multiplicity. *UAI*, 2020.
[17] M. Saerens et al. Adjusting the outputs of a classifier to new a priori probabilities: a simple procedure. *Neural Comput.*, 2002.
[18] M. Sundararajan et al. Axiomatic attribution for deep networks. *ICML*, 2017.
[19] S. Lundberg, S. Lee. A unified approach to interpreting model predictions. *NeurIPS*, 2017.
[20] B. B. Dinh, M. C. Tran, B. Le. Registrations of the follow-up experiments. OSF, 2026. osf.io/vfjtk, osf.io/fr7jv, osf.io/fp8uc